

BACKEND

Problem Statement

Your Challenge: Design and implement a resilient event processing system that can ingest, process, and analyze real-time data streams. Your solution should handle high throughput while maintaining data integrity and providing actionable insights.

- Ingest data from different sources (e.g., sensors, applications, APIs).
- Process the data reliably while handling errors and failures.
- Allow users to access the processed data through a simple API.
- Monitor system performance and detect issues.
- Handle high traffic and keep running smoothly even under stress

Solution approach:

High-Level Architecture

- 1. Event Sources (Producers)**
 - IoT Sensors, Web Applications, APIs, Logs, etc.
 - Producers push events to a message queue (Kafka/Kinesis).
- 2. Event Ingestion Layer**
 - **Kafka/Kinesis Streams** act as a buffer.
 - Supports high throughput & durability.
- 3. Event Processing Layer**
 - **Apache Flink / Apache Spark Streaming**
 - Processes, transforms, enriches, and filters data.
 - Implements error handling (dead-letter queues, retries).
- 4. Storage Layer**
 - **Elasticsearch, DynamoDB, PostgreSQL** (depends on use case).
 - Stores processed data for querying.
- 5. API Layer**
 - **Spring Boot + REST API** for data retrieval.
 - Users can query the processed data.
- 6. Monitoring & Alerting**
 - **Prometheus + Grafana / AWS CloudWatch**
 - Tracks performance, errors, and sends alerts.
- 7. Resilience & Scalability**
 - Auto-scaling groups for microservices.
 - Load testing using **JMeter/K6**.

2. Detailed Breakdown of System Components

1. Event Ingestion Pipeline

- **Tools:** Kafka, AWS Kinesis, RabbitMQ (Kafka preferred).
- **Implementation Steps:**
 - Set up a **Kafka topic** with multiple partitions.
 - Implement **producers** (e.g., Python scripts, Java applications) to send real-time data.
 - Ensure **log retention** and **exactly-once semantics** for data integrity.
- **Key Considerations:**
 - Use **Kafka Connect** for integrating external sources.
 - Enable **idempotent producers** to prevent duplicate events.

2. Detailed Breakdown of System Components

1. Event Ingestion Pipeline

- **Tools:** Kafka.
- **Implementation Steps:**
 - Set up a **Kafka topic** with multiple partitions.
 - Implement **producers** to send real-time data.
 - Ensure **log retention** and **exactly-once semantics** for data integrity.
- **Key Considerations:**
 - Use **Kafka Connect** for integrating external sources.
 - Enable **idempotent producers** to prevent duplicate events.

2. Event Processing with Error Handling

- **Tools:** Apache Flink
- **Implementation Steps:**
 - Consumers **subscribe** to Kafka topics.
 - Process events **in real-time** (transform, enrich, aggregate).
 - Handle **failures** using:
 - **Dead Letter Queues (DLQ)**
 - **Retry mechanisms** (exponential backoff)
 - **Checkpointing & State Management** in Flink.
- **Error Handling Approach:**
 - **Retry up to 3 times**, then push failed events to **DLQ**.
 - Use **Kafka Streams state stores** to keep track of event status.

3. Queryable API for Processed Data

- **Tools:** Spring Boot (Java)
- **Database Options:**
 - **Elasticsearch** (for fast searches)
 - **PostgreSQL** (for structured data, preferred here)
 - **DynamoDB** (for scalable NoSQL storage)
- **API Endpoints:**

Method	Endpoint	Description
GET	/events	Get all processed events
GET	/events/{id}	Get event by ID
GET	/events/all	Get real-time event statistics

4. Monitoring & Alerts

- **Tools:** Prometheus, Grafana, other options can be: AWS CloudWatch, ELK Stack.
- **Key Metrics:**
 - **Event ingestion rate**
 - **Event processing latency**
 - **System errors & failures**
 - **Dead Letter Queue (DLQ) size**
- **Implementation Steps:**
 - Integrate **Prometheus exporters** in Kafka & Flink.
 - Configure **Grafana dashboards** for real-time monitoring.
 - Set up **alerts** for failures using AWS SNS or Slack notifications.

5. System Resilience & Load Testing

- **Resilience Strategies:**
 - **Horizontal Scaling:** Increase Kafka partitions, Flink job managers.
 - **Circuit Breakers:** Use **Resilience4j** to prevent cascading failures.
 - **Retries & Backpressure:** Kafka consumers with exponential backoff.
- **Load Testing:**
 - **Tools:** JMeter
 - **Scenarios:**
 - Simulate **100k+ events/sec**.
 - Test API with **high concurrent requests**.

3. Evaluation Criteria Breakdown

Criteria	How Our System Meets It
Scalability	Kafka partitions, auto-scaling, Flink parallelism
Resilience	DLQs, retries, circuit breakers, checkpointing
Performance	Low-latency processing, high throughput
Design	Modular, microservices-based, clean code
Innovation	Flink for real-time analytics, scalable Kafka-based ingestion

Technology Stack

- **Event Ingestion:** Kafka
- **Processing:** Apache Flink / Apache Spark
- **Storage:** PostgreSQL
- **API:** Spring Boot
- **Monitoring:** Prometheus + Grafana
- **Containerization & Deployment:** Docker

Benefits of DLQ

- **Prevents message loss** even if failures occur.
- **Easier debugging** by storing failed events separately.
- **Improves resilience** by isolating problematic data.

Setup Guide

Install Required Tools

Kafka Setup (Using Docker)

1. Install Docker from [Docker Website](#)
2. Start **Kafka & Zookeeper**: docker-compose up -d
3. Verify Kafka is running: docker ps

Set Up Apache Flink

Install Flink

1. Download Apache Flink:

```
wget https://dlcdn.apache.org/flink/flink-1.15.0-bin-scala_2.12.tgz
```

```
tar -xvzf flink-1.15.0-bin-scala_2.12.tgz
```

```
cd flink-1.15.0
```

1. Start flink

```
./bin/start-cluster.sh
```

Set Up Spring Boot API

◆ Install Java & Maven

1. Install Java 17
2. Install maven

Restart Everything

```
docker-compose down
```

```
docker-compose up -d
```

How to Run the System

1. **Start dependencies**

```
docker-compose up -d
```

1. **Run the Spring Boot Application**

```
mvn spring-boot:run
```

1. **Test Event Processing**

Publish an event

curl -X POST "http://localhost:8080/events?eventType=Temperature&eventData=30C"

Retrieve events

curl http://localhost:8080/events

To check logs

docker-compose logs kafka

Demo images:

The screenshot displays the Postman interface for a REST client. At the top, a list of requests is shown, with the current one being a POST request to `http://localhost:8080/events?eventType=Temperature&eventData=30C`. The main area shows the request configuration with the method set to POST and the URL as `http://localhost:8080/events?eventType=Temperature&eventData=30C`. A blue 'Send' button is visible. Below the URL bar, the 'Params' tab is selected, showing a table of query parameters:

	Key	Value	Bulk Edit
<input checked="" type="checkbox"/>	eventType	Temperature	
<input checked="" type="checkbox"/>	eventData	30C	
	Key	Value	

Below the parameters table, the 'Body' tab is selected, showing the response body in 'Pretty' format. The response is a single line: `1 Event published!`. At the top right of the response area, the status is `200 OK`, the time is `126 ms`, and the size is `180 B`. A 'Save Response' button is also present.